


CS 417 – DISTRIBUTED SYSTEMS

# Week 12: Infrastructure

## High Availability (HA) Clusters



Paul Krzyzanowski

© 2021 Paul Krzyzanowski. No part of this content, may be reproduced or reposted in whole or in part in any manner without the permission of the copyright owner.

# Computer System Design Options

- **Highly Available Systems**

- Incorporate elements of fault-tolerant design
- Component replication, high-quality components
- A fully fault tolerant system will offer non-stop availability ... *but you can't have this!*
- **Problem:**  $\uparrow$  in availability  $\Rightarrow \uparrow$  \$\$

- **High Performance Systems**

- SMP architecture
- Shared memory, shared clock, multiple processors
- **Problems:**
  - Performance gain as  $f(\# \text{ processors})$  is sublinear
  - Contention for resources (bus, memory, devices)
  - The solution is also expensive!

- **Commodity off-the-shelf Systems (COTS)**

- Inexpensive
- Problem: Not reliable and not high performance

# Clustering

Achieve reliability and scalability by interconnecting multiple independent systems

## **Cluster:**

A group of standard, autonomous servers configured so they appear on the network as a single machine

*Single system image*

# Ideally...

- Bunch of off-the shelf machines
- Interconnected on a high-speed LAN
- Appear as one system to users
- Processes are load-balanced across the cluster
  - May migrate
  - May run on different systems
  - All IPC mechanisms and file access available
- Fault tolerant
  - Components may fail
  - Machines may be taken down

# But...

## We don't get all this in off-the-shelf platforms

- Systems design has engineering trade-offs
- Do you need fault-tolerant hardware?
  - Not if your software can work around it
    - Checkpointing, restarting processes, replicated servers, ...
- Do you need high performance?
  - How frequently do processes need to communicate with each other?
    - Scientific computation (e.g., huge matrices) is different from MapReduce or Spark Streaming

# Clustering types

- High availability (HA)
  - Failover cluster
- Supercomputing (HPC)
  - Includes batch processing
- Load balancing (simple form of HA & workload distribution)
- Storage clusters (focus on managing shared storage)

# High Availability (HA) Clustering

# Cluster Components

- Cluster membership
- Heartbeat & heartbeat network
- Quorum
- Configuration & service management
- Storage



# Cluster membership

## Software to manage **cluster membership**

- What are the nodes in the cluster?
- Which nodes in the cluster are currently *alive* (active)?

## We saw this:

- Group Membership Service in virtual synchrony
- GFS master, HDFS NameNode
- Bigtable master
- Pregel master
- MapReduce Master & Spark Cluster Manager

Currently, we most large-scale clusters are custom solutions for specific frameworks. Some components, such as Chubby (Apache Zookeeper) have been adopted by multiple frameworks.

# Quorum

Some members may be dead or disconnected

## Quorum

- Number of elements that must be online for the cluster to function
- Voting algorithm to determine whether the set of nodes has quorum (a majority of nodes to keep running)
- We saw this with Raft consensus (& Paxos): forcing a majority avoids **split-brain**

## Quorum disk

- Shared storage: whichever node can reserve the disk owns it
- Enables systems to resolve who runs a service in small clusters even if the network becomes partitioned

# Types of Quorum

- **Node Majority**

- Each available node can vote
- Need majority (over 50%) of votes for the cluster to continue running
- Best for odd number of nodes, larger clusters

- **Node & Disk Majority** (Microsoft *Disk Witness*)

- Designated shared disk = *disk witness*: counts as a vote
- Need majority of votes to continue running
- Best for an even # of nodes in one site

- **Node & File Share Majority** (Microsoft *File Share Witness*)

- Shared file system = *file share witness* : counts as a vote
- Need majority of votes to continue running
- Windows Server 2019: File Share Witness on USB stick
  - Shared USB storage on router
- Best for an even # of nodes in a multi-site cluster

- **No majority**

- Cluster has quorum if even one node is available and can communicate with a specific disk in the cluster

# Cluster configuration & service management

- **Cluster configuration system & manager**

- UI to manage configuration of systems and software in a cluster
- Administrator has a single point of control

- **Cluster management agent**

- Runs in each cluster node: changes propagate to all nodes
- Tracks cluster membership – removes failed nodes
- Keeps track of quorum – stops cluster when  $\leq$  nodes not active

- **Service management & Scheduler**

- Identify which applications run on which systems
- Specify how failover occurs
  - **Active**: system runs a service
  - **Standby**: Which system(s) can run the service if the active dies
- E.g., MapReduce, Pregel, Spark all use coordinators for their service
- General purpose schedulers: *Google Borg*, *Linux Slurm*

# Disks

# Shared storage access

- If an application can run on any machine, how does it access file data?
- If an application fails over from one machine to another, how does it access its file data?
- Can applications on different machines share files?

# Network (Distributed) File Systems

One option:

- Network file systems: NFS, SMB, AFS, etc.
- Works great for many applications

Concerns

## Availability

- Address with replication (most file systems offer little)

## Performance

- Remote systems on a LAN vs. local bus access
- Overhead of remote operating system & network stack
- Point of congestion
- Look at GFS/HDFS to distribute file data across lots of servers  
... or other parallel file systems, such as Lustre, GlusterFS, or Ceph

# Shared disks & Cluster file systems

## Shared disk

- Allows multiple systems to share access to disk drives
- Works well if there isn't much contention
  - ... but you can't have multiple systems reading/writing/caching the same disk blocks

## Cluster File System

- Client runs a file system accessing a shared disk at the **block level**
  - *vs. a distributed file system, which access at a file-system level*
- No client/server roles, no disconnected modes
- All nodes are peers and access a shared disk(s)
- **Distributed Lock Manager (DLM)**
  - Process to ensure mutual exclusion for disk access
  - Provides inode-based locking and caching control
  - Not needed for local file systems on a shared disk



# Cluster File Systems

## Examples:

- IBM General Parallel File System (GPFS)
- Microsoft Cluster Shared Volumes (CSV)
- Oracle Cluster File System (OCFS)
- Red Hat Global File System (GFS2)

## Linux GFS2 (no relation to Google GFS)

- Cluster file system accessing storage at a **block level**
- **Cluster Logical Volume Manager (CLVM)**:  
Volume management of cluster storage
- **Global Network Block Device (GNBD)**:  
Block level storage access over ethernet: cheap way to access block-level storage

# The alternative: shared nothing

## Shared nothing

- No shared devices
- Each system has its own storage resources
- No need to deal with DLMS
- If a machine A needs resources on B, A sends a message to B
  - If B fails, storage requests have to be switched over to a live node

Requires **exclusive** access to shared storage

Rely on active replication of changes or ...

- Multiple nodes may have access to shared storage
- Only one node is granted exclusive access at a time – *one owner*
- Exclusive access changed on failover

# SAN: Computer-Disk interconnect

## SAN = Storage Area Network

- Separate network between nodes and storage arrays
  - Fibre channel
  - iSCSI
- Any node can be configured to access any storage through a fibre channel switch

## Acronyms

- **DAS**: Direct Attached Storage
- **SAN**: block-level access to a disk via a network
- **NAS**: file-level access to a remote file system (NFS, SMB, ...)

# Failover

# HA issues

- How do you detect failover?
- How long does it take to detect?
- How does a dead application move/restart?
- Where does it move to?

# Heartbeat network

- Machines need to detect faulty systems
  - **Heartbeat**: Periodic “ping” mechanism
  - An “are you alive” message
- Need to distinguish ***system faults*** from ***network faults***
  - Useful to maintain redundant networks
  - Avoid split-brain issues in systems without quorum (e.g., a 2-node cluster)
- Once you know who is dead or alive, then determine a course of action

# Failover Configuration Models

- **Active/Passive**
  - Requests go to active system
  - Passive nodes do nothing until they're needed
  - Passive nodes maintain replicated state (e.g., SMR/Virtual Synchrony)
  - Example: Chubby
- **Active/Active**
  - Any node can handle a request
  - Failed workload goes to remaining nodes
  - Replication must be  $N$ -way for  $N$  active nodes
  - Example: GFS chunks
- **Active/Passive:  $N+M$** 
  - $M$  dedicated failover node(s) for  $N$  active nodes

# Design options for failover

- **Cold failover**

- Application restart
- *Example: map and reduce workers in MapReduce*

- **Warm failover**

- Restart last checkpointed image
- Relies on application checkpointing itself periodically
- *Example: Pregel*

- **Hot failover**

- Application state is synchronized across systems
  - E.g., replicated state machines or lockstep synchronization at the CPU level
- Spare is ready to run immediately
- May be difficult at a fine granularity, prone to software faults (e.g., what if a specific set of inputs caused the software to die?)
- *Example: Chubby*



# Design options for failover

With either type of failover ...

## **Multi-directional failover**

- Failed applications migrate to or restart on available systems

*And possibly*

## **Cascading failover**

- If the backup system fails, application can be restarted on another surviving system

# IP Address Takeover (IPAT)

Depending on the deployment:

- **Ignore**

- IP addresses of services don't matter. A load balancer, name server, or coordinator will identify the correct machine

- **Take over IP address**

- A node in an active/passive configuration may need to take over the IP address of a failed node

- **Take over MAC address**

- MAC address takeover may be needed if we cannot guarantee that other nodes will flush their ARP cache

- **Listen on multiple addresses**

- A node in an active/active configuration may need to listen on multiple IP addresses

# Hardware support for High Availability

- **Hot-pluggable components**
  - Minimize downtime for component swapping
  - E.g., disks, power supplies, CPU/memory boards
- **Redundant devices**
  - Redundant power supplies
  - Parity on memory
  - Mirroring on disks (or RAID for HA)
  - Switchover of failed components
- **Diagnostics**
  - On-line identification & service

# Fencing

- **Fencing**: method of isolating a node from a cluster
  - Apply to failed node
  - Disconnect I/O to ensure data integrity
  - Avoid problems with Byzantine failures
  - Avoids problems with *fail-restart*
    - Restarted node has not kept up to date with state changes
- **Types of fencing**
  - **Power fencing**: shut power off a node
  - **SAN fencing**: disable a Fibre Channel port to a node
  - **System service fencing**: disable access to a global network block device (GNBD) server
  - **Software fencing**: remove server processes from the group
    - E.g., virtual synchrony

# Cluster software hierarchy

## Example: Windows Server cluster abstractions

### Top tier: Cluster abstractions

- Failover manager (what needs to be started/restarted?)
- Resource monitor (what's going on?)
- Cluster registry (who belongs in the cluster?)

### Middle tier: Distributed operations

- Global status update
- Membership
- Quorum (and leader election)

### Bottom tier: OS and drivers

- Cluster disk driver, cluster network drivers
- IP address takeover

# The End